


# Learning curves — ONLINE TRAINING ROLLOUTS, not evaluation (mean $\pm$ SE across 8 seeds)



TRAINING ROLLOUTS (exploring policy, resetting from the TRAINING bank) — NOT the evaluation protocol.  
Endpoints are NOT the table numbers: eval uses mode/sampled actions on the pole / rnd30 start sets, final checkpoint only.  
Do not compare against PHASE\_RESULTS \$2.y / \$3 / \$5.1 or the HRI\_results tables.