


Learning curves, WORLD MODEL (r2dreamer) only — ONLINE TRAINING ROLLOUTS, not evaluation (mean $\pm$ SE across 8 seeds)



WORLD MODEL (r2dreamer) ONLY — RLPD persists no training-time metrics and DP is offline (no online rollouts), so neither can produce this curve without new compute; the per-learner comparison exists for FINAL numbers only. TRAINING ROLLOUTS (exploring policy, resetting from the TRAINING bank) — NOT the evaluation protocol. Endpoints are NOT the table numbers: eval uses mode/sampled actions on the pole / rnd30 start sets, final checkpoint only. Do not compare against PHASE_RESULTS \$2.y / \$3 / \$5.1 or the HRI_results tables.