

{dv3} local -- dv3 (DreamerV3 losses in the r2dreamer chassis), pixels

observation: pixels; ladder: nested\_sparse10; human n=4, machine n=4; budget 1M; TRAINING RECORD, sampled actions, policy's own training starts; line = mean over seeds, band = 95 % bootstrap CI (2000 resamples), truncated where fewer than 3 seeds have data. Datasets: H = human demonstrations (dHfull\_all, 74 tapes), M = machine (DDPfull\_first, 72 tapes from a Diffusion Policy teacher), P = planner (planner72, 72 motion-planner tapes), T = r2dreamer teacher (72 tapes from one human-trained pixel r2dreamer). Hardware: pop-os workstation, AMD 5950X (AVX2) + local GPU.

