

{r2dreamer} cluster -- r2dreamer (the port's contrastive representation loss), pixels observation: pixels; ladder: nested_sparse10; human n=16, machine n=16, planner72 n=8, r2teacher n=8; budget 1M/2M; TRAINING RECORD, sampled actions, policy's own training starts; line = mean over seeds, band = 95 % bootstrap CI (2000 resamples), truncated where fewer than 3 seeds have data. Datasets: H = human demonstrations (dHfull_all, 74 tapes), M = machine (dDPfull_first, 72 tapes from a Diffusion Policy teacher), P = planner (planner72, 72 motion-planner tapes), T = r2dreamer teacher (72 tapes from one human-trained pixel r2dreamer).

