



{RLPD} cluster -- RLPD (DrQ-style shared encoder, LN critic ensemble), pixels

observation: pixels; ladder: nested_sparse10; human n=9, machine n=8, planner72 n=8, r2teacher n=5; budget 1M;
 TRAINING RECORD, sampled actions, policy's own training starts; line = mean over seeds, band = 95 % bootstrap CI (2000
 resamples), truncated where fewer than 3 seeds have data. Datasets: H = human demonstrations (dHfull_all, 74 tapes),
 M = machine (dDPfull_first, 72 tapes from a Diffusion Policy teacher), P = planner (planner72, 72 motion-planner
 tapes), T = r2dreamer teacher (72 tapes from one human-trained pixel r2dreamer). x-axis = decisions x action_repeat 4
 = sim frames; farside/slide_event/home are ABSENT from the RLPD episode record (absent, not zero).


picked




placed (v2)




farside



slide_event



home



tipped

