

{RLPD} cluster -- RLPD (DrQ-style shared encoder, LN critic ensemble), pixels
 observation: pixels; ladder: nested_sparse10; human n=9, machine n=8, planner72 n=8, r2teacher n=1; budget 1M;
 TRAINING RECORD, sampled actions, policy's own training starts; line = mean over seeds, band = 95 % bootstrap CI (2000
 resamples), truncated where fewer than 3 seeds have data; datasets with fewer than 3 seeds (r2dreamer teacher) are
 drawn one thin line per seed, no band. Datasets: H = human demonstrations (dHfull_all, 74 tapes), M = machine
 (dDPfull_first, 72 tapes from a Diffusion Policy teacher), P = planner (planner72, 72 motion-planner tapes), T =
 r2dreamer teacher (72 tapes from one human-trained pixel r2dreamer). x-axis = decisions x action_repeat 4 = sim
 frames; farside/slide_event/home are ABSENT from the RLPD episode record (absent, not zero).

