


{RLPD} cluster -- `home` TRAINING CURVE derived from ep_rew_mean/10 (pixels, nested_sparse10)



{RLPD} (DrQ-style encoder, LN critic ensemble, pixels, ladder nested_sparse10). The RLPD episode record has no `home` flag, so this curve is DERIVED from the Slurm log: SB3's `ep\_rew\_mean` is the mean return over the last 100 online episodes, and the ladder pays only +10 at `home` (tip penalty 0), so $\text{ep_rew_mean} / 10 = \text{the rolling-100 `home` rate}$. It is a rolling-100 window, not the rolling-30 window of the other learning curves; sampled actions, the policy's own training starts. Thin lines = seeds; thick line + band = mean and 95 % bootstrap CI (2000 resamples) where ≥ 3 seeds have data. Records past 250k decisions (13 (ag) seeds trained on) are cut. Every seed with a log is drawn, including human s8, which is outside the $n=8$ design of the results table. Datasets: H = human demonstrations (dHfull_all, 74 tapes), M = machine (dDPfull_first, 72 tapes from a Diffusion Policy teacher), P = planner (planner72, 72 motion-planner tapes), T = r2dreamer teacher (72 tapes from one human-trained pixel r2dreamer).