




Rise-time LAG relative to `picked`. For each run, $\text{lag}(\text{phase}) = \text{rise}(\text{rate} \geq 0.1, \text{phase}) - \text{rise}(\text{rate} \geq 0.1, \text{'picked'})$, where rise is the first online sim step at which the run's rolling-30-episode rate for that flag reaches the threshold.

TRAINING RECORD: sampled actions on the policy's own training starts; observation pixels, ladder nested_sparse10. {dv3} = DreamerV3 losses in the r2dreamer chassis, {r2dreamer} = the port's contrastive representation loss. Demonstrations: human = dHfull_all (74 tapes), machine = dDPfull_first (72 tapes). One point per seed (jittered), thick bar = median over the seeds that crossed, the small number under each group is how many crossed. A seed whose phase never reaches the threshold inside its record is an OPEN marker on the dotted top edge and is counted there; it is never plotted as lag 0 and never enters the median. A seed whose `picked` itself never crosses has no lag at all and is reported in the panel title. {RLPD} is absent from this figure: its episode record carries no `farside`/`slide\_event`/`home` flag (absent, not zero). The nested_ramp control is absent: it pays a different ladder. Budgets 1 M (ae/af) or 2 M (ag) online sim steps; the {dv3} local seeds ran on the pop-os AVX2 GPU workstation, the cluster conditions on pax GPU nodes.