




Rise-time LAG relative to 'picked'. For each run, $\text{lag}(\text{phase}) = \text{rise}(\text{rate} \geq 0.5, \text{phase}) - \text{rise}(\text{rate} \geq 0.5, \text{'picked'})$, where rise is the first online sim step at which the run's rolling-30-episode rate for that flag reaches the threshold.

TRAINING RECORD: sampled actions on the policy's own training starts; observation pixels, ladder nested_sparse10. {dv3} = DreamerV3 losses in the r2dreamer chassis, {r2dreamer} = the port's contrastive representation loss. Datasets: H = human demonstrations (dHfull_all, 74 tapes), M = machine (dDPfull_first, 72 tapes from a Diffusion Policy teacher), P = planner (planner72, 72 motion-planner tapes), T = r2dreamer teacher (72 tapes from one human-trained pixel r2dreamer). One point per seed (jittered), thick bar = median over the seeds that crossed, the small number under each group is how many crossed. A seed whose phase never reaches the threshold inside its record is an OPEN marker on the dotted top edge and is counted there; it is never plotted as lag 0 and never enters the median; a run still inside its budget that has not crossed is PENDING, not plotted, and counted in the panel title. A seed whose 'picked' itself never crosses has no lag at all and is reported in the panel title. {RLPD} is absent from this figure: its episode record carries no 'farside'/'slide_event'/'home' flag (absent, not zero). The nested_ramp control is absent: it pays a different ladder. Budgets 1 M (ae/af) or 2 M (ag) online sim steps; the {dv3} local seeds ran on the pop-os AVX2 GPU workstation, the cluster conditions on pax GPU nodes.