



(a) TRAINING RECORD -- mean `home` over the last 20 % of each run's online steps (sampled actions, the policy's own training starts); one point per seed, bar = arm mean, n above each column



(b) EVAL CELLS -- mean over every rnd30 MODE milestone cell at ≥ 0.5 M online steps (deterministic actions, 30 fixed random starts, fresh process); one point per seed, bar = arm mean of `home`



● home ▲ picked

Steady-state performance. {dv3} = DreamerV3 losses, {r2dreamer} = the port's contrastive loss, {RLPD} = DrQ-style encoder in the LN critic ensemble; observation pixels, ladder nested_sparse10. Ticks: dv3 loc / dv3 clu / r2d clu / RLPD / ramp ctl, H = human demonstrations (dHfull_all, 74 tapes), M = machine (dDPfull_first, 72 tapes). Budgets: 1 M (ae/af) or 2 M (ag) online sim steps; {RLPD} 250k decisions = 1 M sim frames. {dv3} loc ran on the pop-os AVX2 GPU workstation, every other condition on the pax cluster. `ramp ctl' uses nested\_ramp, not nested\_sparse10. {RLPD} logs no `home' flag in its episode record, so panel (a) has no RLPD column (absent, not zero). In (a) a HOLLOW point is a run that has not finished its budget: its last 20 % is not a steady state. A seed with no eval cell at ≥ 0.5 M is absent from (b), not zero.