



(a) TRAINING RECORD -- mean `home` over the last 20 % of each run's online steps (sampled actions, the policy's own training starts); one point per seed, bar = arm mean, n above each column



(b) EVAL CELLS -- world models: mean over every rnd30 MODE milestone cell at  $\geq 0.5$  M online steps; {RLPD}: the 250k rnd30 MODE cell; {Diffusion Policy}: the rnd30 SAMPLE cell (30 fixed random starts, fresh process); one point per seed, bar = dataset mean of `home`



Steady-state performance. {dv3} = DreamerV3 losses, {r2dreamer} = the port's contrastive loss, {RLPD} = DrQ-style encoder in the LN critic ensemble, {Diffusion Policy} = lerobot DP (eval only, sampled actions, no episode record so no point in (a)); observation pixels, ladder nested\_sparse10. Ticks: dv3 loc / dv3 clu / r2d clu / RLPD / DP / ramp ctl; H = human demonstrations (dHfull\_all, 74 tapes), M = machine (dDPfull\_first, 72 tapes from a Diffusion Policy teacher), P = planner (planner72, 72 motion-planner tapes), T = r2dreamer teacher (72 tapes from one human-trained pixel r2dreamer). Budgets: 1 M (ae/af) or 2 M (ag) online sim steps; {RLPD} 250k decisions = 1 M sim frames. {dv3} loc ran on the pop-os AVX2 GPU workstation, every other condition on the pax cluster. `ramp ctl' uses nested\_ramp, not nested\_sparse10. {RLPD} logs no `home' flag in its episode record, so panel (a) has no RLPD column (absent, not zero). In (a) a HOLLOW point is a run that has not finished its budget: its last 20 % is not a steady state. A seed with no eval cell at  $\geq 0.5$  M is absent from (b), not zero.